

QuantOptima: Rethinking Neural Network Codebooks

An Independent Research Project on the Mathematical Limits of
Weight Compression

Andrei Mihai Pană
National College of Computer Science “Tudor Vianu”

Abstract

Large Language Models and deep computer vision architectures use billions of parameters. Storing them in standard 32-bit floating point takes too much memory space, making them impossible to run on normal hardware. To fix this, I compressed the memory arrays to 8-bit or 4-bit representations. This experiment investigates if standard industry formats (like INT8 and FP8) are optimally designed, or if calculating a custom, dynamic codebook using a Lloyd-Max k-means algorithm provides better accuracy. I found that minimizing the theoretical rounding error (MSE) does not necessarily improve real model accuracy, which suggests the neural network relies on critical outliers in the weight distribution.

Version 2 Improvements and Methodological Revisions

This second version of the research report introduces several critical methodological improvements to ensure the validity of the conclusions:

- **Statistical Rigor:** Previous accuracy differences between formats fell within the margin of error. All benchmarks are now run across 5 different random seeds, reporting the mean and standard deviation. Key comparisons are backed by McNemar's statistical test to distinguish real performance gains from statistical noise.
- **Format Equalization:** The FP8 encoder was upgraded to properly support subnormal numbers (E4M3 specification). Additionally, all formats (including FP8 and Posit8) now receive identical per-tensor scaling, ensuring a perfectly fair comparison without hardcoded advantages.
- **Code-Report Alignment:** The INT8 implementation was updated to be strictly symmetric (where 0 maps exactly to 0), perfectly aligning the codebase with the mathematics described in the report.
- **Expanded Tournament:** Standard 4-bit formats (INT4 and FP4) were added to the benchmark to provide a complete picture of extreme compression at lower bit-widths.
- **Rigorous Crossover Analysis:** The Kurtosis experiment was completely redesigned. Synthetic data is no longer artificially clamped, allowing the true outlier tails to form naturally. Furthermore, the live activations of the network were extracted and analyzed to explain the discrepancy between weight compression and activation compression.

1. Introduction to the Memory Bottleneck

When a neural network is trained, it learns patterns by adjusting the values of its weights. In frameworks like PyTorch, these weights are stored as `float32` variables. A single `float32` takes 4 bytes of memory.

If you take a model like GPT-3, which has 175 billion parameters, just storing the weights requires about 700 Gigabytes of VRAM. A high-end consumer graphics card like the RTX 4090 only has 24GB. This massive gap is why quantization is not just a fun trick, but a strictly mandatory step for running modern AI.

Quantization means forcing those continuous 32-bit decimals into a smaller bucket of discrete values. If I use 8 bits, I only have $2^8 = 256$ unique values

to represent my numbers. If I drop to 4 bits, I only get $2^4 = 16$ unique values. The core question of this whole project is: **If I only have 16 slots, what are the exact mathematically perfect values I should pick to fill them?**

2. Architecture and Baseline Training

Before I can compress anything, I need a real model with real weights. I built a Convolutional Neural Network (CNN) from scratch to classify images from the CIFAR-10 dataset.

The architecture of CifarCNN is designed with three spatial convolution blocks followed by a dense linear classifier:

- conv1: Expands the 3 RGB channels to 32 channels.
- conv2: Expands 32 channels to 64 channels.
- conv3: Expands 64 channels to 128 channels.
- fc1: A fully connected linear layer containing over 500,000 weights.
- fc2: The final output layer mapping to the 10 image classes.

I trained this network on my CPU for 5 epochs using the Adam optimizer with a learning rate of 0.001. The Adam optimizer dynamically adjusts the step size for every single weight using momentum, which helps the network converge fast.

After training across 5 random seeds, the baseline models achieved a mean accuracy of **75.48% $\pm$ 0.53%** on the unseen test dataset.

To see what I am actually trying to compress, I extracted the raw memory tensors from the trained model and plotted their histograms.




Figure 1: The weight distribution for the first convolutional layer.




Figure 2: The weight distribution for the massive fully connected layer.

As shown in the figures above, the distributions for layers like conv1 and fc1 naturally clump really tight around zero. They look basically like a standard bell curve. Most of the network's knowledge is concentrated in numbers very close to 0.0. However, notice the tiny, almost invisible tails stretching outwards to the extremes on the X-axis.

These exact tails are the “outliers”. They make quantization extremely hard, because if you stretch your 256 available values to cover the far outliers, you lose precision in the middle where the bulk of the weights are.

3. Memory & Hardware Compression Ratios

Before looking at the accuracy, I measured exactly how much physical memory my formats saved. The CifarCNN model contains exactly 617, 226 parameters.

Precision Format	Total Weights Memory	Compression Ratio	Hardware Status
FP32 (Baseline)	2.46 MB	1.0x	Heavy
INT8 / FP8	617 KB	4.0x	Optimized
NF4 / INT4	308 KB	8.0x	Ultra-Light

Table 1: Memory footprint analysis of the trained CifarCNN model.

By dropping to 4-bit, I reduced the model size by 800%. In a real-world scenario with a 70-Billion parameter language model, this is the exact difference between needing a cluster of server GPUs versus running it entirely on a single gaming laptop.

4. Benchmarking Standard Formats (Phase 1)

To see how the industry handles the accuracy drop, I wrote custom bitwise encoders and decoders for four major quantization formats from scratch, without relying on external compression libraries.

- **INT8 (Integer 8-bit)** Uses a rigid, evenly spaced grid. I implemented a symmetric version, which finds the absolute maximum value in the tensor and divides it by 127 to create a scale factor s . Every weight is then divided by s and rounded to the nearest integer.

$$Q(x) = \frac{x}{s} * s \quad (1)$$

It snapped the weights into 256 evenly spaced buckets. Really simple but effective.

- **FP8 (Floating Point 8-bit, E4M3)** Instead of a flat grid, this format uses a 1-bit sign, a 4-bit exponent, and a 3-bit mantissa. Because exponents grow exponentially, this format naturally clusters a lot of unique values very close to zero, and spreads them out further apart as the numbers get bigger.

- **Posit8 (Posit Standard)** A highly advanced alternative to floating point. It uses “regime bits” (a run-length encoding sequence) to dynamically shift precision. If a number is near 1.0 or 0.0, Posit steals bits from the exponent to give to the mantissa, providing extreme sub-decimal precision right where neural networks need it most.
- **NF4 (NormalFloat 4-bit)** This format doesn’t use math formulas like exponents. It assumes the neural network weights follow a perfect Gaussian bell curve. It calculates the exact quantiles of a standard normal distribution and hardcodes those 16 specific values as the codebook.

I applied these formats to my trained CNN by overriding the memory arrays in place.

Format Engine	Bit-Width	Mean MSE	Test Accuracy (Mean $\pm$ Std)
FP32 (Baseline)	32-bit	0.0000	75.48% $\pm$ 0.53%
INT8 (Symmetric)	8-bit	< 0.0001	75.48% $\pm$ 0.52%
FP8 (E4M3)	8-bit	< 0.0001	75.44% $\pm$ 0.51%
Posit8 (es=0)	8-bit	0.0008	71.97% $\pm$ 1.34%
NF4	4-bit	0.0007	74.66% $\pm$ 0.53%
INT4 (Symmetric)	4-bit	0.0014	73.64% $\pm$ 0.42%
FP4 (E2M1)	4-bit	0.0012	72.73% $\pm$ 0.66%

Table 2: Phase 1 & 2 Standard Formats Benchmark Results on CIFAR-10 (over 5 seeds).

Following the fixes to subnormal handling and the introduction of fair per-tensor scaling, the previous “Accuracy Paradox” where INT8 magically outperformed FP8 disappeared. Both 8-bit formats perform statistically identically to the FP32 baseline.

Answering the Verification Question: In Version 1 of this report, the FP32 baseline scored 75.05% and INT8 scored 75.07%. Out of the 10,000 images in the CIFAR-10 test dataset, a difference of 0.02 percentage points translates to exactly **2 images**. My entire hypothesis that “INT8 acts as a mathematical regularizer” was based on a random variance of just two images. The McNemar test across multiple seeds now mathematically proves this difference was nothing but statistical noise (p-value: 0.7699 between INT8 and FP8). There is no “mathematical regularizer” effect—at 8 bits, they are simply equivalent to the

baseline. When dropping to 4-bit formats, we observe that the standard INT4 format outperforms the exponential FP4 format on this CNN architecture.

5. Building the “Perfect” Codebook (Phase 2)

The standard formats are just educated guesses. NF4 guesses the weights are a perfect bell curve. INT8 guesses an even spread is fine. But what if we don’t guess?

Instead of assuming the shape of the weights, I wrote a Lloyd-Max algorithm to find the theoretically perfect 1D codebook for every single layer independently. This is basically a 1D k-means clustering algorithm.

The algorithm has three steps:

1. Initialize 256 random values from the array as the starting “centroids”.
2. Assign every single float in the weight matrix to its closest centroid.
3. Update the centroid’s position to be the exact mean of all the points assigned to it.
4. Repeat until the centroids stop moving.

To make this run fast on millions of weights without slow Python loops, I implemented the math purely with vectorized PyTorch tensor operations, using `unsqueeze` for matrix broadcasting distances and `bincount` to accumulate the cluster sums instantly.

The goal of this algorithm is to strictly minimize the Mean Squared Error (MSE) between the original weights vector $\vec{W}$ and the quantized weights vector $\vec{Q}$.

$$\text{MSE} = \frac{1}{N} \sum_{i=1}^N (W_i - Q_i)^2 \quad (2)$$

Algebraic Proof of Convergence: I can prove that moving the centroid to the mean (average) is the mathematically perfect way to minimize the error without using any calculus limits or derivatives.

Let S_k be the set of all weights assigned to a specific cluster. Let C be the centroid value we are trying to find. Let μ be the true arithmetic mean of the weights in S_k :

$$\mu = \frac{\sum_{x \in S_k} x}{|S_k|} \quad (1)$$

We want to minimize the sum of the squared distances for this cluster:

$$\text{Error} = \sum_{x \in S_k} (x - C)^2 \quad (1)$$

By using normal algebra, we can add and subtract μ inside the square:

$$\text{Error} = \sum_{x \in S_k} (x - \mu + \mu - C)^2 \quad (1)$$

Expanding the quadratic expression $(a + b)^2 = a^2 + 2ab + b^2$:

$$\text{Error} = \sum_{x \in S_k} [(x - \mu)^2 + 2(x - \mu)(\mu - C) + (\mu - C)^2] \quad (1)$$

We can split this into three separate sums. The middle sum has a constant $2(\mu - C)$ which we can pull out:

$$\text{Error} = \sum_{x \in S_k} (x - \mu)^2 + 2(\mu - C) \sum_{x \in S_k} (x - \mu) + \sum_{x \in S_k} (\mu - C)^2 \quad (1)$$

Look at the middle term: $\sum (x - \mu)$. The sum of the differences between a set of numbers and their own mean is always exactly zero. So the entire middle term is deleted. Because $(\mu - C)^2$ is a constant, summing it $|S_k|$ times is just multiplying it.

The equation simplifies to:

$$\text{Error} = \sum_{x \in S_k} (x - \mu)^2 + |S_k|(\mu - C)^2 \quad (1)$$

The first term is just the natural variance of the data, we cannot change it. The only part we can control is the second term. Since a squared number is always positive, the absolute lowest value we can make the second term is 0. This only happens if:

$$C = \mu \quad (1)$$

This algebraically proves that updating the centroid to the mean always perfectly minimizes the MSE! I ran this “perfect” math engine on my network. Because it optimizes specifically for each layer, it achieved the lowest possible MSE in the world. But look at the final results in the table below.

Quantization Strategy	Bits	Mean MSE	Final Accuracy (Mean $\pm$ Std)
INT8 (Symmetric)	8-bit	< 0.0001	$75.48\% \pm 0.52\%$
OPTIM (Lloyd-Max)	8-bit	< 0.0001	$75.41\% \pm 0.45\%$
NF4 (Normal Distribution)	4-bit	0.0007	$74.66\% \pm 0.53\%$
OPTIM (Lloyd-Max)	4-bit	0.0004	$74.31\% \pm 0.46\%$

Table 3: OPTIM achieves lower MSE, but often underperforms standard formats in test accuracy.

Why did the MSE metric fail to predict accuracy? Lloyd-Max places the codebook values where the massive clump of weights is, completely ignoring the rare extreme outliers because there are not enough of them to pull the mean. The results are consistent with the hypothesis that the neural network relies on those extreme outliers to function correctly. Standard MSE does not weight these outliers heavily enough.

6. The Star Graph: The Entire Project in One Image

To summarize the whole research question visually, I plotted the accuracy of every format against its bit-width constraint.




Figure 3: Accuracy vs. Bits. The ultimate showdown between standard and optimal codebooks.

This “Star Graph” illustrates the performance gap. At 8-bit, the standard INT8 performs similarly to the FP32 baseline, while the unweighted OPTIM line falls underneath. When dropping to 4-bit, the gap widens again, with the hardcoded NF4 bell-curve outperforming the dynamic OPTIM codebook. It suggests that searching for lower unweighted MSE may be the wrong direction for model compression.

7. The Kurtosis Crossover (Phase 3)

To find exactly when formats fail, I needed to isolate the shape of the data. I wrote a script to generate synthetic distributions where I could manually control the excess kurtosis (κ).

Kurtosis measures how “fat” the tails of a distribution are.

$$\kappa = \frac{1}{\sigma^4} \left(\frac{1}{N} \sum_{i=1}^N (x_i - \mu)^4 \right) - 3 \quad (1)$$

I blended a Uniform distribution (no tails, flat block) with a Cauchy distribution (extreme, infinite tails) using an interpolation parameter α :

$$D_{\text{mixed}} = (1 - \alpha) * D_{\text{uniform}} + \alpha * D_{\text{cauchy}} \quad (1)$$

I measured the MSE of INT8, FP8, and Posit8 as the data slowly grew fatter tails.




Figure 4: Format Supremacy Inversion as kurtosis increases.

I evaluated the inversion point. As I slide the distribution to the right (higher kurtosis), the standard INT8 format sees a significant increase in MSE compared to FP8.

Because INT8 is a rigid grid, an extreme outlier forces the absolute maximum boundary to stretch to wide. When the grid stretches to cover the massive outlier, it crushes all the important middle values into the zero bucket. FP8 and Posit8 handle this much better because their exponents allow them to reach far outliers without sacrificing the center resolution.

The Activations Bonus: The crossover analysis above was strictly performed on the model’s static weights. However, real-world data passing through the network behaves differently. By hooking into the PyTorch forward pass, I extracted the live activations for the first 5 batches of the test set. While the weights for layers like `fc1` have a kurtosis of 8.11, the **activations** flowing into that exact same layer have a staggering kurtosis of 21.73. Activations in deeper layers like `fc2` reached a kurtosis of 23.10. This perfectly aligns with current literature (like `LLM.int8()`): while network weights are relatively well-behaved and can be compressed effectively with standard integer formats, the activations are riddled with massive, unpredictable outliers that heavily favor floating-point formats like FP8.




Figure 5: Activation histograms for conv1 (Kurtosis 2.17) and fc1 (Kurtosis 21.73). Notice the extreme outliers stretching far out on the X-axis in the fc1 activations, requiring FP8’s dynamic range.

8. Conclusion

This project completely changed how I view data compression. I discovered that minimizing MSE is a trap. A mathematically “perfect” codebook fails in the real world because it smooths out the chaotic outliers that the neural network relies on to make predictions. Standard formats like INT8 and NF4

are actually already operating at near-maximum efficiency for their respective distributions.

9. Limitations

While this experiment successfully explores the relationship between MSE and quantization accuracy, there are several strict limitations to my methodology that must be acknowledged:

- **Scale of the Network:** I trained a small CNN with 600,000 parameters on CIFAR-10. This is fundamentally different from a 70-Billion parameter Transformer. Large Language Models (LLMs) have massive activation outliers that CNNs simply do not generate. The kurtosis crossover point might happen much earlier in an LLM.
- **Post-Training Quantization (PTQ) Only:** I only applied the codebooks **after** the network was fully trained. I did not test Quantization-Aware Training (QAT), where the network learns to actively route its logic around the grid limitations while it is still training.
- **Hardware Simulation vs. Real Execution:** My custom Python encoders fake the quantization by casting the low-bit representations back to FP32 for the PyTorch forward pass. I did not write custom CUDA kernels to measure real-world inference speedups, so I can only prove memory compression, not execution latency.